

CASE STUDY-1

1. Introduction

1.1 Business Problem

In the present system, the e-commerce platform will send a feedback mail to customers after the product is delivered. The customers can give ratings out of 5, also can write down some comments/reviews about the product that he/she has purchased. Using these reviews and ratings, e-commerce platform will rate the products, which helps other people to get the insights about the quality of the product. But according to seller perspective, these reviews will play crucial role to improve the business. But many times, customers would not give any ratings or reviews. How to predict the review score that a customer could give? This is the problem in e-commerce business. Also, the problem can be extended as "Is it possible to predict the review rating that a customer could give before actually he gives the rating?". If this problem is solved, then it is also possible to predict the rating for which customer had not given any rating. In this case study, my objective is to try to solve this problem, that is Predicting the e-commerce customer satisfaction.

For this case study, I have taken the dataset given by Olist, which is an e-commerce platform in Brazil. Olist connects small businesses all over Brazil to customers with a single contract. Olist has provided over 100k order information that were placed between 2016 to 2018. Similar to all other e-commerce platforms, Olist also send feedback form to customers after the estimated delivery date to get the reviews and ratings. Now, Olist wants to improve the business as well as provide the better service to customers by using the customer satisfaction information. For that it needs to predict the review ratings before the user will give actual ratings. So, my approach is to address this business problem using data science, which is a scientific way to solve this business problem.

1.2 ML formulation of business problem:

To solve the business problem using data science, it is needed to pose that problem as classical machine learning problem. First of all, since the data has target variable, it is supervised ML problem. Further we need to predict the satisfaction of customers, that is predicting the ratings. Ratings are discrete ranging from 1 to 5. Hence it is a multi-class classification problem. We have 5 class labels, So, we can treat the problem as 5-class classification ML problem. Our goal is to predict the rating before the user give the rating/review. Hence, we should not consider data regarding review message, comments, etc as features.

1.3 Business constraints:

- There is no strict low latency requirement. But model should not take too much time for predicting, Since we should get the prediction before the user give.
- Low ratings like 1,2,3 are very important with respect to business improvement. So, misclassification of low ratings would cost loss of customers. Hence misclassifications are crucial.
- Since low ratings are crucial, if we get the interpretations of the output, that will be better.

1.4 Performance Metrics: (possible)

- Multi-class confusion matrix
- Micro F1 score
- Precision Recall curve for each class
- Multi-class log loss
- Balanced accuracy score

1.5 Research-Papers/Solutions/Architectures/Kernels

- Existing solution: <https://www.kaggle.com/andresionek/predicting-customer-satisfaction> (<https://www.kaggle.com/andresionek/predicting-customer-satisfaction>)
- Repeat buyer prediction for e-commerce <https://www.kdd.org/kdd2016/papers/files/adf0160-liuA.pdf> (<https://www.kdd.org/kdd2016/papers/files/adf0160-liuA.pdf>)
- LightGBM Classifier: <https://papers.nips.cc/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf> (<https://papers.nips.cc/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf>)
- SMOTE technique: <https://arxiv.org/pdf/1106.1813.pdf> (<https://arxiv.org/pdf/1106.1813.pdf>)
- Blog on undersampling, oversampling, SMOTE, ensemble models: <https://xang1234.github.io/louvain/> (<https://xang1234.github.io/louvain/>)
- https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html (https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html)
- <https://machinelearningmastery.com/random-oversampling-and-undersampling-for-imbalanced-classification/> (<https://machinelearningmastery.com/random-oversampling-and-undersampling-for-imbalanced-classification/>)
- <https://towardsdatascience.com/using-data-science-to-predict-negative-customer-reviews-2abbdfbf3d82> (<https://towardsdatascience.com/using-data-science-to-predict-negative-customer-reviews-2abbdfbf3d82>)

In []:

2. Exploratory Data Analysis

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import folium

import datetime
```

2.1 Loading the data files

In [2]:

```
#loading all csv files
customer      = pd.read_csv("olist_customers_dataset.csv")
geo_location  = pd.read_csv("olist_geolocation_dataset.csv")
items         = pd.read_csv("olist_order_items_dataset.csv")
payments      = pd.read_csv("olist_order_payments_dataset.csv")
reviews       = pd.read_csv("olist_order_reviews_dataset.csv")
orders        = pd.read_csv("olist_orders_dataset.csv")
products      = pd.read_csv("olist_products_dataset.csv")
sellers       = pd.read_csv("olist_sellers_dataset.csv")
translation   = pd.read_csv("product_category_name_translation.csv")
```

In [3]:

```
def overview(dataframe):
    """This function will return the overview of the dataframe"""

    print("Shape of the dataframe is : {}".format(dataframe.shape))
    print("***30")
    print("Information about features : ",dataframe.info())
    print("***30")
    print("Total number of null values : \n ",dataframe.isnull().sum())
    print("***30")

    return dataframe.head(3)
```

2.2 Overview of all data

```
overview(customer)
```

```
Shape of the dataframe is : (99441, 5)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99441 entries, 0 to 99440
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   customer_id                          99441 non-null  object
1   customer_unique_id                   99441 non-null  object
2   customer_zip_code_prefix             99441 non-null  int64
3   customer_city                        99441 non-null  object
4   customer_state                       99441 non-null  object
dtypes: int64(1), object(4)
memory usage: 3.8+ MB
Information about features : None
*****
Total number of null values :
    customer_id          0
customer_unique_id      0
customer_zip_code_prefix 0
customer_city           0
customer_state          0
dtype: int64
*****
```

Out[4]:

	customer_id	customer_unique_id	customer_zip_code_prefix	customer_city	customer_s
0	06b8999e2fba1a1fbc88172c00ba8bc7	861eff4711a542e4b93843c6dd7febb0	14409	franca	
1	18955e83d337fd6b2def6b18a428ac77	290c77bc529b7ac935b93aa66c333dc3	9790	sao bernardo do campo	
2	4e7b3e00288586ebd08712fdd0374a03	060e732b5b29e8181a18229c7b0b2b5e	1151	sao paulo	

```
overview(geo_location)
```

```
Shape of the dataframe is : (1000163, 5)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000163 entries, 0 to 1000162
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   geolocation_zip_code_prefix          1000163 non-null  int64
1   geolocation_lat                      1000163 non-null  float64
2   geolocation_lng                      1000163 non-null  float64
3   geolocation_city                    1000163 non-null  object
4   geolocation_state                   1000163 non-null  object
dtypes: float64(2), int64(1), object(2)
memory usage: 38.2+ MB
Information about features : None
*****
Total number of null values :
    geolocation_zip_code_prefix    0
    geolocation_lat                 0
    geolocation_lng                 0
    geolocation_city                0
    geolocation_state               0
    dtype: int64
*****
```

Out[5]:

	geolocation_zip_code_prefix	geolocation_lat	geolocation_lng	geolocation_city	geolocation_state
0	1037	-23.545621	-46.639292	sao paulo	SP
1	1046	-23.546081	-46.644820	sao paulo	SP
2	1046	-23.546129	-46.642951	sao paulo	SP

In [6]:

overview(items)

Shape of the dataframe is : (112650, 7)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 112650 entries, 0 to 112649
Data columns (total 7 columns):
Column Non-Null Count Dtype
--- ----- -----
0 order_id 112650 non-null object
1 order_item_id 112650 non-null int64
2 product_id 112650 non-null object
3 seller_id 112650 non-null object
4 shipping_limit_date 112650 non-null object
5 price 112650 non-null float64
6 freight_value 112650 non-null float64
dtypes: float64(2), int64(1), object(4)
memory usage: 6.0+ MB
Information about features : None

Total number of null values :
 order_id 0
order_item_id 0
product_id 0
seller_id 0
shipping_limit_date 0
price 0
freight_value 0
dtype: int64

Out[6]:

	order_id	order_item_id	product_id	seller_id	st
0	00010242fe8c5a6d1ba2dd792cb16214	1	4244733e06e7ecb4970a6e2683c13e61	48436dade18ac8b2bce089ec2a041202	2f
1	00018f77f2f0320c557190d7a144bdd3	1	e5f2d52b802189ee658865ca93d83a8f	dd7ddc04e1b6c2c614352b383efe2d36	2f
2	000229ec398224ef6ca0657da4fc703e	1	c777355d18b72b67abbeef9df44fd0fd	5b51032eddd242adc84c38acab88f23d	2f

In [7]:

overview(payments)

Shape of the dataframe is : (103886, 5)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103886 entries, 0 to 103885
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- ----- -----
0 order_id 103886 non-null object
1 payment_sequential 103886 non-null int64
2 payment_type 103886 non-null object
3 payment_installments 103886 non-null int64
4 payment_value 103886 non-null float64
dtypes: float64(1), int64(2), object(2)
memory usage: 4.0+ MB
Information about features : None

Total number of null values :
 order_id 0
payment_sequential 0
payment_type 0
payment_installments 0
payment_value 0
dtype: int64

Out[7]:

	order_id	payment_sequential	payment_type	payment_installments	payment_value
0	b81ef226f3fe1789b1e8b2acac839d17	1	credit_card	8	99.33
1	a9810da82917af2d9aefd1278f1dcfa0	1	credit_card	1	24.39
2	25e8ea4e93396b6fa0d3dd708e76c1bd	1	credit_card	1	65.71

In [8]:

overview(reviews)

Shape of the dataframe is : (100000, 7)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 7 columns):
Column Non-Null Count Dtype
--- ----- -
0 review_id 100000 non-null object
1 order_id 100000 non-null object
2 review_score 100000 non-null int64
3 review_comment_title 11715 non-null object
4 review_comment_message 41753 non-null object
5 review_creation_date 100000 non-null object
6 review_answer_timestamp 100000 non-null object
dtypes: int64(1), object(6)
memory usage: 5.3+ MB
Information about features : None

Total number of null values :
 review_id 0
order_id 0
review_score 0
review_comment_title 88285
review_comment_message 58247
review_creation_date 0
review_answer_timestamp 0
dtype: int64

Out[8]:

	review_id		order_id	review_score	review_comment_title	review_comment_
0	7bc2406110b926393aa56f80a40eba40	73fc7af87114b39712e6da79b0a377eb		4		NaN
1	80e641a11e56f04c1ad469d5645fdfde	a548910a1c6147796b98fdf73dbeba33		5		NaN
2	228ce5500dc1d8e020d8d1322874b6f0	f9e4b658b201a9f2ecdecbb34bed034b		5		NaN

In [9]: overview(orders)

```
Shape of the dataframe is : (99441, 8)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99441 entries, 0 to 99440
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   order_id                             99441 non-null  object
1   customer_id                           99441 non-null  object
2   order_status                           99441 non-null  object
3   order_purchase_timestamp               99441 non-null  object
4   order_approved_at                      99281 non-null  object
5   order_delivered_carrier_date           97658 non-null  object
6   order_delivered_customer_date          96476 non-null  object
7   order_estimated_delivery_date          99441 non-null  object
dtypes: object(8)
memory usage: 6.1+ MB
Information about features : None
*****
Total number of null values :
  order_id                0
customer_id              0
order_status             0
order_purchase_timestamp 0
order_approved_at        160
order_delivered_carrier_date 1783
order_delivered_customer_date 2965
order_estimated_delivery_date 0
dtype: int64
*****
```

Out[9]:

	order_id	customer_id	order_status	order_purchase_timestamp	order_approv
0	e481f51cbdc54678b7cc49136f2d6af7	9ef432eb6251297304e76186b10a928d	delivered	2017-10-02 10:56:33	201711
1	53cdb2fc8bc7dce0b6741e2150273451	b0830fb4747a6c6d20dea0b8c802d7ef	delivered	2018-07-24 20:41:37	201803
2	47770eb9100c2d0c44946d9cf07ec65d	41ce2a54c0b03bf3443c3d931a367089	delivered	2018-08-08 08:38:49	201808

```
In [10]: overview(products)
```

```
Shape of the dataframe is : (32951, 9)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32951 entries, 0 to 32950
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   product_id                            32951 non-null  object
1   product_category_name                 32341 non-null  object
2   product_name_lenght                  32341 non-null  float64
3   product_description_lenght           32341 non-null  float64
4   product_photos_qty                   32341 non-null  float64
5   product_weight_g                     32949 non-null  float64
6   product_length_cm                    32949 non-null  float64
7   product_height_cm                    32949 non-null  float64
8   product_width_cm                     32949 non-null  float64
dtypes: float64(7), object(2)
memory usage: 2.3+ MB
Information about features : None
*****
Total number of null values :
    product_id                0
product_category_name        610
product_name_lenght          610
product_description_lenght    610
product_photos_qty           610
product_weight_g              2
product_length_cm             2
product_height_cm             2
product_width_cm              2
dtype: int64
*****
```

Out[10]:

	product_id	product_category_name	product_name_lenght	product_description_lenght	product_photos
0	1e9e8ef04dbcff4541ed26657ea517e5	perfumaria	40.0	287.0	
1	3aa071139cb16b67ca9e5dea641aaa2f	artes	44.0	276.0	
2	96bd76ec8810374ed1b65e291975717f	esporte_lazer	46.0	250.0	

```
In [11]: overview(sellers)
```

```
Shape of the dataframe is : (3095, 4)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3095 entries, 0 to 3094
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   seller_id                            3095 non-null  object
1   seller_zip_code_prefix               3095 non-null  int64
2   seller_city                          3095 non-null  object
3   seller_state                         3095 non-null  object
dtypes: int64(1), object(3)
memory usage: 96.8+ KB
Information about features : None
*****
Total number of null values :
    seller_id                0
seller_zip_code_prefix      0
seller_city                 0
seller_state                0
dtype: int64
*****
```

Out[11]:

	seller_id	seller_zip_code_prefix	seller_city	seller_state
0	3442f8959a84dea7ee197c632cb2df15	13023	campinas	SP
1	d1b65fc7debc3361ea86b5f14c68d2e2	13844	mogi guacu	SP
2	ce3ad9de960102d0677a81f5d0bb7b2d	20031	rio de janeiro	RJ

```
In [12]: overview(translation)

Shape of the dataframe is : (71, 2)
*****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 71 entries, 0 to 70
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   product_category_name                 71 non-null    object
1   product_category_name_english         71 non-null    object
dtypes: object(2)
memory usage: 1.2+ KB
Information about features : None
*****
Total number of null values :
  product_category_name          0
product_category_name_english    0
dtype: int64
*****
```

```
Out[12]:
```

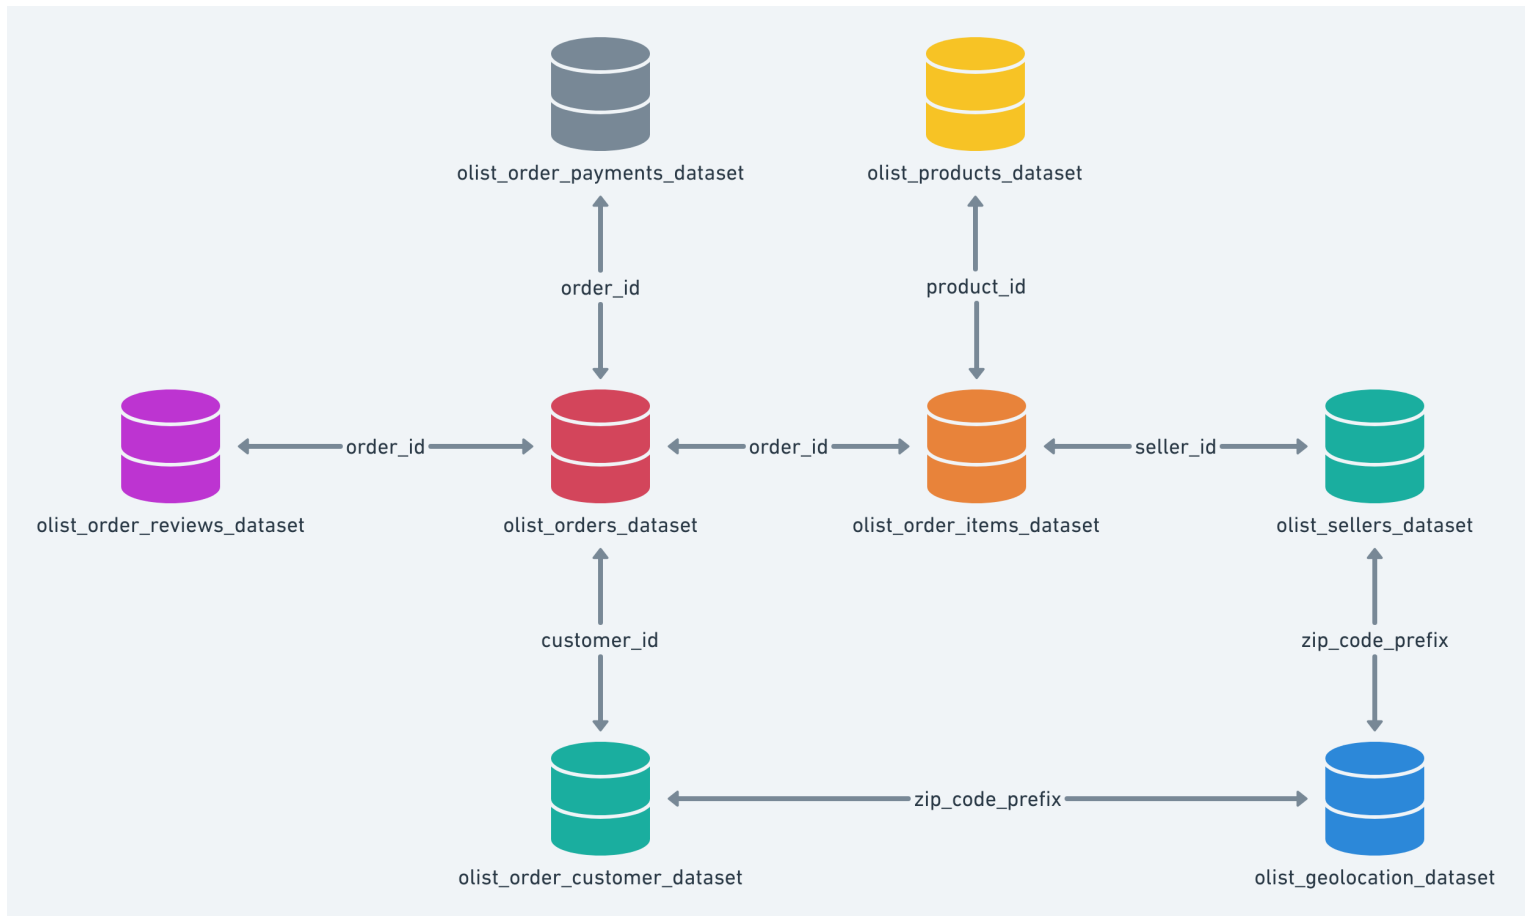
	product_category_name	product_category_name_english
0	beleza_saude	health_beauty
1	informatica_acessorios	computers_accessories
2	automotivo	auto

Observations:

- * There are null values present in reviews data, orders data, product data. Rest of the dataframes do not have any null values.
- * Translation data has the product category names in english, i.e. english version of product categories. So we can use this to get the product category names in english.
- * geo_location data has 1000163 data points, which seems very large compare to the data points of other dataframes. There is possibility of duplicate points present in geo_locations.
- * To get the complete informations, we need to merge these sub dataframes to get single dataframe. Merging should be done according to the data schema that is given below.

2.3 Merge the dataframes

Data schema



In [13]: *#zip_code_prefix column has different names in customer,geo_location,sellers dataframes
#Change these names to zip_code_prefix*

```
customer.rename(columns = {"customer_zip_code_prefix":"zip_code_prefix"},inplace=True)
geo_location.rename(columns = {"geolocation_zip_code_prefix":"zip_code_prefix"},inplace=True)
sellers.rename(columns = {"seller_zip_code_prefix":"zip_code_prefix"},inplace=True)
```

```
In [14]: #drop the duplicate values in geo_location (zip_code_prefix)

print("Number of rows before dropping duplicates in geo_location : ",geo_location.shape[0])
geo_location.drop_duplicates(subset="zip_code_prefix",keep="first",inplace=True)
print("Number of rows after dropping duplicates in geo_locations : ",geo_location.shape[0])

Number of rows before dropping duplicates in geo_location : 1000163
Number of rows after dropping duplicates in geo_locations : 19015
```

In []:

```
In [15]: #customer merging with geo_location (left join to preserve customer info)
geo_customer      = pd.merge(customer,geo_location,how="left",on="zip_code_prefix")

#sellers merging with geo_locations (left join to preserve sellers info)
geo_seller        = pd.merge(sellers,geo_location,how="left",on="zip_code_prefix")

#payment merging with order
payment_order     = pd.merge(payments,orders,on="order_id")

#product merging with item
product_item      = pd.merge(products,items,on="product_id")

#payment_order merging with reviews
pay_order_review  = pd.merge(payment_order,reviews,on="order_id")

#pay_order_review merging with geo_customer
review_customer   = pd.merge(pay_order_review,geo_customer,on="customer_id")

#prod_item merging with geo_seller
prod_item_seller  = pd.merge(product_item,geo_seller,on="seller_id")

#Finally review_customer with prod_item_seller
data              = pd.merge(review_customer,prod_item_seller,on="order_id",suffixes=("_customer","_seller"))
```

```
In [16]: num_tot_rows = data.shape[0]
print("Shape of the final dataframe (data) is : ",data.shape)

Shape of the final dataframe (data) is : (118315, 47)
```

In []:

2.4 Data Cleaning

```
In [17]: print("All the column names in the data : \n",data.columns)
```

```
All the column names in the data :
Index(['order_id', 'payment_sequential', 'payment_type',
       'payment_installments', 'payment_value', 'customer_id', 'order_status',
       'order_purchase_timestamp', 'order_approved_at',
       'order_delivered_carrier_date', 'order_delivered_customer_date',
       'order_estimated_delivery_date', 'review_id', 'review_score',
       'review_comment_title', 'review_comment_message',
       'review_creation_date', 'review_answer_timestamp', 'customer_unique_id',
       'zip_code_prefix_customer', 'customer_city', 'customer_state',
       'geolocation_lat_customer', 'geolocation_lng_customer',
       'geolocation_city_customer', 'geolocation_state_customer', 'product_id',
       'product_category_name', 'product_name_lenght',
       'product_description_lenght', 'product_photos_qty', 'product_weight_g',
       'product_length_cm', 'product_height_cm', 'product_width_cm',
       'order_item_id', 'shipping_limit_date', 'price',
       'freight_value', 'zip_code_prefix_seller', 'seller_city',
       'seller_state', 'geolocation_lat_seller', 'geolocation_lng_seller',
       'geolocation_city_seller', 'geolocation_state_seller'],
      dtype='object')
```

We should only use the informations which are upto the product delivery. We don't use the review text, comment,review creation date like features for the prediction. Hence it is better to drop these columns now.

```
In [18]: data.drop(['review_id', 'review_comment_title', 'review_comment_message',  
                'review_creation_date', 'review_answer_timestamp'], axis=1, inplace=True)
```

```
In [19]: #Checking for nul values  
data.isnull().sum()
```

```
Out[19]: order_id                                0  
payment_sequential                             0  
payment_type                                   0  
payment_installments                           0  
payment_value                                  0  
customer_id                                    0  
order_status                                   0  
order_purchase_timestamp                       0  
order_approved_at                             15  
order_delivered_carrier_date                   1254  
order_delivered_customer_date                 2588  
order_estimated_delivery_date                 0  
review_score                                   0  
customer_unique_id                             0  
zip_code_prefix_customer                      0  
customer_city                                 0  
customer_state                                 0  
geolocation_lat_customer                     317  
geolocation_lng_customer                     317  
geolocation_city_customer                    317  
geolocation_state_customer                   317  
product_id                                    0  
product_category_name                       1709  
product_name_lenght                          1709  
product_description_lenght                   1709  
product_photos_qty                           1709  
product_weight_g                             20  
product_length_cm                            20  
product_height_cm                            20  
product_width_cm                             20  
order_item_id                                 0  
seller_id                                     0  
shipping_limit_date                           0  
price                                          0  
freight_value                                0  
zip_code_prefix_seller                       0  
seller_city                                   0  
seller_state                                  0  
geolocation_lat_seller                       265  
geolocation_lng_seller                       265  
geolocation_city_seller                      265  
geolocation_state_seller                     265  
dtype: int64
```

```
In [20]: data["order_status"].value_counts()
```

```
Out[20]: delivered      115728  
shipped                1255  
canceled               570  
invoiced               376  
processing             376  
unavailable            7  
approved              3  
Name: order_status, dtype: int64
```

```
In [21]: not_delivered = data[data["order_status"]!="delivered"]["order_status"].value_counts().sum()  
  
print("Total number of orders which are not delivered : ", not_delivered)  
print("percentage of orders which are not delivered : ", (not_delivered*100/data.shape[0]).round(3), "%"  
)
```

```
Total number of orders which are not delivered : 2587  
percentage of orders which are not delivered : 2.187 %
```

```
In [22]: #We are interested in products that are delivered.
#We don't need the products which are not delivered yet/ cancelled. hence Let's consider only delivered orders.
# Also there are only 2.2% of data which are not delivered.
```

```
data = data[data["order_status"]=="delivered"]
```

```
In [23]: #Since we cannot impute datetime, Lets drop the rows which has null values in order_approved_at,order_d
elivered_carrier_date,order_delivered_customer_date
#These columns has very little number of null values.
data.dropna(subset=["order_approved_at","order_delivered_carrier_date","order_delivered_customer_date"
],axis=0,inplace=True)
```

```
In [24]: data.isnull().sum()
```

```
Out[24]: order_id                                0
payment_sequential                             0
payment_type                                  0
payment_installments                          0
payment_value                                 0
customer_id                                   0
order_status                                  0
order_purchase_timestamp                     0
order_approved_at                           0
order_delivered_carrier_date                 0
order_delivered_customer_date                0
order_estimated_delivery_date                0
review_score                                0
customer_unique_id                           0
zip_code_prefix_customer                     0
customer_city                                0
customer_state                               0
geolocation_lat_customer                     303
geolocation_lng_customer                     303
geolocation_city_customer                    303
geolocation_state_customer                   303
product_id                                   0
product_category_name                        1637
product_name_lenght                          1637
product_description_lenght                   1637
product_photos_qty                           1637
product_weight_g                             20
product_length_cm                           20
product_height_cm                           20
product_width_cm                            20
order_item_id                                0
seller_id                                    0
shipping_limit_date                          0
price                                         0
freight_value                                0
zip_code_prefix_seller                       0
seller_city                                  0
seller_state                                 0
geolocation_lat_seller                       261
geolocation_lng_seller                       261
geolocation_city_seller                      261
geolocation_state_seller                     261
dtype: int64
```

```
In [ ]:
```

There are 261 rows in which seller location informations are null. Similarly 303 rows in which customer locations are null. We cannot impute these values, also these numbers are small, Hence we can drop these.

```
In [25]: data.dropna(subset=["geolocation_lng_seller","geolocation_state_customer"],axis=0,inplace=True)
```

```
In [26]: data.isnull().sum()
```

```
Out[26]: order_id                                0
payment_sequential                             0
payment_type                                  0
payment_installments                          0
payment_value                                 0
customer_id                                   0
order_status                                  0
order_purchase_timestamp                     0
order_approved_at                           0
order_delivered_carrier_date                 0
order_delivered_customer_date                0
order_estimated_delivery_date                0
review_score                                 0
customer_unique_id                           0
zip_code_prefix_customer                     0
customer_city                                0
customer_state                               0
geolocation_lat_customer                     0
geolocation_lng_customer                     0
geolocation_city_customer                    0
geolocation_state_customer                   0
product_id                                   0
product_category_name                        1631
product_name_lenght                          1631
product_description_lenght                   1631
product_photos_qty                           1631
product_weight_g                             20
product_length_cm                            20
product_height_cm                            20
product_width_cm                             20
order_item_id                                0
seller_id                                    0
shipping_limit_date                          0
price                                         0
freight_value                                0
zip_code_prefix_seller                       0
seller_city                                  0
seller_state                                 0
geolocation_lat_seller                       0
geolocation_lng_seller                       0
geolocation_city_seller                      0
geolocation_state_seller                     0
dtype: int64
```

```
In [ ]:
```

```
In [27]: data.dropna(subset=["product_photos_qty", "product_width_cm"], axis=0, inplace=True)
```

```
In [28]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 113509 entries, 0 to 118314
Data columns (total 42 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   order_id                                  113509 non-null  object
1   payment_sequential                       113509 non-null  int64
2   payment_type                             113509 non-null  object
3   payment_installments                    113509 non-null  int64
4   payment_value                           113509 non-null  float64
5   customer_id                             113509 non-null  object
6   order_status                            113509 non-null  object
7   order_purchase_timestamp                113509 non-null  object
8   order_approved_at                      113509 non-null  object
9   order_delivered_carrier_date            113509 non-null  object
10  order_delivered_customer_date            113509 non-null  object
11  order_estimated_delivery_date            113509 non-null  object
12  review_score                             113509 non-null  int64
13  customer_unique_id                       113509 non-null  object
14  zip_code_prefix_customer                113509 non-null  int64
15  customer_city                           113509 non-null  object
16  customer_state                           113509 non-null  object
17  geolocation_lat_customer                 113509 non-null  float64
18  geolocation_lng_customer                 113509 non-null  float64
19  geolocation_city_customer                113509 non-null  object
20  geolocation_state_customer                113509 non-null  object
21  product_id                              113509 non-null  object
22  product_category_name                    113509 non-null  object
23  product_name_lenght                     113509 non-null  float64
24  product_description_lenght               113509 non-null  float64
25  product_photos_qty                       113509 non-null  float64
26  product_weight_g                         113509 non-null  float64
27  product_length_cm                       113509 non-null  float64
28  product_height_cm                       113509 non-null  float64
29  product_width_cm                        113509 non-null  float64
30  order_item_id                            113509 non-null  int64
31  seller_id                                113509 non-null  object
32  shipping_limit_date                      113509 non-null  object
33  price                                    113509 non-null  float64
34  freight_value                           113509 non-null  float64
35  zip_code_prefix_seller                  113509 non-null  int64
36  seller_city                             113509 non-null  object
37  seller_state                             113509 non-null  object
38  geolocation_lat_seller                   113509 non-null  float64
39  geolocation_lng_seller                   113509 non-null  float64
40  geolocation_city_seller                  113509 non-null  object
41  geolocation_state_seller                 113509 non-null  object
dtypes: float64(14), int64(6), object(22)
memory usage: 37.2+ MB
```

```
In [29]: data.shape
```

```
Out[29]: (113509, 42)
```

```
In [30]: # % of data that we have drooped
print("Percentage of data that we have dropped is : ",round((num_tot_rows-data.shape[0])*100/num_tot_rows,3),"%")
```

```
Percentage of data that we have dropped is : 4.062 %
```

So we have lost roughly 4% of data, which is not bad. Because we have 96% of data which don't have nul values.

```
In [31]: data[data.duplicated()==True]
```

Out[31]:

	order_id	payment_sequential	payment_type	payment_installments	payment_value	
4	ba78997921bbcdc1373bb41e913ab953	1	credit_card	8	107.78	7a5d8efa
66	c0db7d31ace61fc360a3eaa34dd3457c	1	credit_card	5	65.71	80c0276f
74	95442deb81a5d91c97c0df96b431634a	1	boleto	1	368.98	daddb54f
75	95442deb81a5d91c97c0df96b431634a	1	boleto	1	368.98	daddb54f
188	14d9794bbb53614d12cc2df6a045f82b	1	boleto	1	116.70	666a4eb7f
...
116252	b9eb9fbcca54dfce1d087c752d0f4b35	1	boleto	1	35.01	38243190
117353	f7c5c5ff5045e13c98901bbb8e871d4	1	credit_card	10	231.27	1d5f730
117453	169ab175fb915582d84c0c5c95bb0fe3	1	credit_card	10	741.10	fe02528a
117624	03515a836bb855b03f7df9dee520a8fc	1	credit_card	10	106.38	153e6d88f
117883	c88b1d1b157a9999ce368f218a407141	1	credit_card	4	42.77	ae0fb7bf

401 rows × 42 columns

```
In [32]: #remove the duplicates
data.drop_duplicates(keep="first",inplace=True)
```

```
In [33]: data.shape
```

Out[33]: (113108, 42)

```
In [34]: #lets change the datetime features to correct format
data["order_purchase_timestamp"] = pd.to_datetime(data["order_purchase_timestamp"])
data["order_approved_at"] = pd.to_datetime(data["order_approved_at"])
data["order_delivered_carrier_date"] = pd.to_datetime(data["order_delivered_carrier_date"])
data["order_delivered_customer_date"] = pd.to_datetime(data["order_delivered_customer_date"])
data["order_estimated_delivery_date"] = pd.to_datetime(data["order_estimated_delivery_date"])
data["shipping_limit_date"] = pd.to_datetime(data["shipping_limit_date"])
```

I found some problem with seller_city, seller_state information in sellers dataframe.

In [35]: data[data["geolocation_state_seller"]!=data["seller_state"]]

Out[35]:

		order_id	payment_sequential	payment_type	payment_installments	payment_value	
53	1dcf0c8cd36ffaf57784fdbc90079310		1	credit_card	3	157.15	34955e04f
238	591083bc42b589c7052118aa83118e76		5	voucher	1	20.00	276df8efc
240	591083bc42b589c7052118aa83118e76		3	voucher	1	20.00	276df8efc
242	591083bc42b589c7052118aa83118e76		2	voucher	1	20.00	276df8efc
244	591083bc42b589c7052118aa83118e76		6	voucher	1	15.21	276df8efc
...
117074	d3037ef55ee7f4d6c892da7493f4912f		1	credit_card	5	154.47	a92c7dc3
117345	15cd24278cb6a373aa6bb2ca34837e16		1	boleto	1	24.60	61a5483
117599	cab69d0a36811c4945aec2d4f744c8b5		1	credit_card	10	230.19	2468eac8
117658	652dd239023471248957a1a2d5173b60		1	credit_card	3	140.44	35cc89fc7
117984	6e9583ad9400699e091e33534fa84aa4		1	credit_card	3	236.29	89fe9585e

653 rows × 42 columns

There are 653 record in which seller_state and geo_location_state_seller are not matching.
So, let us look deeply about this issue.

In [36]: geo_location[(geo_location["geolocation_state"]=="SC") & (geo_location["geolocation_city"]=="itajai")]

Out[36]:

	zip_code_prefix	geolocation_lat	geolocation_lng	geolocation_city	geolocation_state
909293	88310	-26.891052	-48.701608	itajai	SC
909296	88309	-26.895500	-48.687952	itajai	SC
909297	88306	-26.935823	-48.627747	itajai	SC
909300	88311	-26.877642	-48.705441	itajai	SC
909307	88302	-26.920172	-48.657852	itajai	SC
909318	88317	-26.876740	-48.750403	itajai	SC
909319	88308	-26.928509	-48.700390	itajai	SC
909322	88303	-26.912429	-48.677381	itajai	SC
909332	88304	-26.899380	-48.679729	itajai	SC
909374	88312	-26.908691	-48.703615	itajai	SC
909405	88316	-26.945308	-48.718467	itajai	SC
910010	88318	-27.037284	-48.861505	itajai	SC
912083	88313	-26.971308	-48.683232	itajai	SC

In [37]: geo_location[(geo_location["geolocation_state"]=="SP") & (geo_location["geolocation_city"]=="itajai")]

Out[37]:

	zip_code_prefix	geolocation_lat	geolocation_lng	geolocation_city	geolocation_state
--	-----------------	-----------------	-----------------	------------------	-------------------


```
In [38]: sellers[(sellers["seller_state"]=="SP") & (sellers["seller_city"]=="itajai")]
```

Out[38]:

	seller_id	zip_code_prefix	seller_city	seller_state
1303	0bae85eb84b9fb3bd773911e89288d54	88301	itajai	SP
1449	52b53f7061969fe471d119b6195da864	88301	itajai	SP

In []:

So, I googled the actual city of itajai, Brazil. Actually state is SC (Santa Catarina). Not SP (São Paulo).

So, sellers data contains wrong information about this.

geo_location is accurate.

Hence it is better to drop seller_state, seller_city, since we have the correct information in geolocation data.i.e.

geolocation_state_seller, geo_location_city_seller.


Since seller data has some wrong info, let us check customer data also.

```
In [39]: data[data["geolocation_state_customer"]!=data["customer_state"]]
```

Out[39]:

order_id	payment_sequential	payment_type	payment_installments	payment_value	customer_id	order_status	order_purchase
----------	--------------------	--------------	----------------------	---------------	-------------	--------------	----------------

0 rows × 42 columns



```
In [40]: data[data["geolocation_city_customer"]!=data["customer_city"]][["geolocation_city_customer","customer_c
ity"]]
```

Out[40]:

	geolocation_city_customer	customer_city
57	jaguariúna	jaguariuna
58	guaraí	guarai
59	guaraí	guarai
60	guaraí	guarai
76	são lourenço	sao lourenco
...
118147	picuí	picui
118157	brasilíia	brasilíia
118250	são paulo	sao paulo
118273	maringá	maringa
118291	são paulo	sao paulo

8254 rows × 2 columns

We can see that customer_state and geolocation_customer_state are same.

But in city names, customer_city names are in english. But geolocation_city_customer has different notation. But names are same.

Here we shall drop customer_city, since geolocation_city_customer has same notations as all seller_city has.

```
In [41]: data.drop(["seller_state","seller_city","customer_city","customer_state"],axis=1,inplace=True)

data.rename(columns =
             {"geolocation_state_seller":"seller_state","geolocation_city_seller":"seller_city",
              "geolocation_lng_seller":"lng_seller","geolocation_lat_seller":"lat_seller",
              "geolocation_lng_customer":"lng_customer","geolocation_lat_customer":"lat_customer"
             },
             {"geolocation_city_customer":"customer_city","geolocation_state_customer":"customer
state"},
             ,inplace=True)
```

There are 3 records in the data, which has 0 installment. This is weird. Also these 3 records present in rating 5.
Let us drop these 3 records with installment = 0.

```
In [42]: data[data["payment_installments"]==0]
```

```
Out[42]:
```

	order_id	payment_sequential	payment_type	payment_installments	payment_value
56288	744bade1fc9ff3f31d860ace076d422	2	credit_card	0	58.69
91637	1a57108394169c0b47d8f876acc9ba2d	2	credit_card	0	129.94
91638	1a57108394169c0b47d8f876acc9ba2d	2	credit_card	0	129.94

3 rows × 38 columns

```
In [43]: data.drop(data.index[[53769,87577,87578]],axis=0,inplace=True)
```

```
In [44]: data[data["payment_installments"]==0]
```

```
Out[44]:
```

order_id	payment_sequential	payment_type	payment_installments	payment_value	customer_id	order_status	order_purchase
----------	--------------------	--------------	----------------------	---------------	-------------	--------------	----------------

0 rows × 38 columns

map the product_category translation to data

```
In [45]: data = pd.merge(data,translation,how="left",on="product_category_name")
```

```
In [46]: data[data["product_category_name_english"].isnull()==True]["product_category_name"].value_counts()
```

```
Out[46]: portateis_cozinha_e_preparadores_de_alimentos    14
pc_gamer    9
Name: product_category_name, dtype: int64
```

```
In [47]: translation[translation["product_category_name"]=="portateis_cozinha_e_preparadores_de_alimentos"]
```

```
Out[47]:
```

product_category_name	product_category_name_english
-----------------------	-------------------------------

Translation data deos not has portateis_cozinha_e_preparadores_de_alimentos, pc_gamer.
portateis_cozinha_e_preparadores_de_alimentos english translation is kitchen_laptops_and_food_preparators
pc_gamer is pc_gamer.

Let us fill the null values in english translation by these values.

```
In [48]: null_1 = data[data["product_category_name"]=="portateis_cozinha_e_preparadores_de_alimentos"]["product_category_name_english"]
null_2 = data[data["product_category_name"]=="pc_gamer"]["product_category_name_english"]
```

```
In [49]: data.loc[null_1.index,"product_category_name_english"] = "kitchen_laptops_and_food_preparators"
data.loc[null_2.index,"product_category_name_english"] = "pc_gamer"
```

```
In [50]: data.drop("product_category_name",inplace=True,axis=1)
data.rename(columns={"product_category_name_english":"product_category_name"},inplace=True)
```

```
In [ ]:
```

```
In [51]: #Let us save the data as final_data.csv, so that we do not need to run the above cells again.
data.to_csv("final_data.csv")
```

In []:

Exploratory Data Analysis

In [52]:

```
#Load the data
data = pd.read_csv("final_data.csv")
data.drop("Unnamed: 0",axis=1,inplace=True)
```

In [53]:

```
#info about the data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 113105 entries, 0 to 113104
Data columns (total 38 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   order_id                             113105 non-null object
1   payment_sequential                   113105 non-null int64
2   payment_type                         113105 non-null object
3   payment_installments                113105 non-null int64
4   payment_value                       113105 non-null float64
5   customer_id                         113105 non-null object
6   order_status                        113105 non-null object
7   order_purchase_timestamp            113105 non-null object
8   order_approved_at                   113105 non-null object
9   order_delivered_carrier_date        113105 non-null object
10  order_delivered_customer_date       113105 non-null object
11  order_estimated_delivery_date       113105 non-null object
12  review_score                        113105 non-null int64
13  customer_unique_id                  113105 non-null object
14  zip_code_prefix_customer            113105 non-null int64
15  lat_customer                        113105 non-null float64
16  lng_customer                        113105 non-null float64
17  customer_city                       113105 non-null object
18  customer_state                      113105 non-null object
19  product_id                          113105 non-null object
20  product_name_lenght                113105 non-null float64
21  product_description_lenght         113105 non-null float64
22  product_photos_qty                 113105 non-null float64
23  product_weight_g                   113105 non-null float64
24  product_length_cm                  113105 non-null float64
25  product_height_cm                  113105 non-null float64
26  product_width_cm                   113105 non-null float64
27  order_item_id                      113105 non-null int64
28  seller_id                          113105 non-null object
29  shipping_limit_date                113105 non-null object
30  price                              113105 non-null float64
31  freight_value                      113105 non-null float64
32  zip_code_prefix_seller              113105 non-null int64
33  lat_seller                          113105 non-null float64
34  lng_seller                          113105 non-null float64
35  seller_city                         113105 non-null object
36  seller_state                       113105 non-null object
37  product_category_name               113105 non-null object
dtypes: float64(14), int64(6), object(18)
memory usage: 32.8+ MB
```


Let us check the distribution of class label to understand the class balance

```
In [54]: # checking the balance of the data
print("Counts :\n ",data["review_score"].value_counts())
print("***30)
print("Percentage distribution :\n ",100*data["review_score"].value_counts(normalize=True))
print("***30)

plt.bar(data["review_score"].value_counts().index,data["review_score"].value_counts(),color="darkblue")

plt.xlabel("Ratings")
plt.ylabel("Counts")
plt.title("Distribution of review scores")
plt.show()
```

```
Counts :
5    64637
4    21725
1    13241
3     9573
2     3929
Name: review_score, dtype: int64
*****
Percentage distribution :
5    57.147783
4    19.207816
1    11.706821
3     8.463817
2     3.473763
Name: review_score, dtype: float64
*****
```



We can see that review scores distributed in **J** like shape. This is the typical distribution of e-commerce review scores. Very high number of rating 5 followed by 4 and then rating 1, rating 2 and 3 are less compared to other ratings.

The dataset is highly imbalanced.


There are only 3.5% datapoints which belongs to rating 2.

About 57% of datapoints belongs to rating 5.

Correlation between features

In [55]: *#Reference documentation: https://seaborn.pydata.org/examples/many_pairwise_correlations.html*

```
correlation = data.corr()
#mask the above triangle
mask = np.triu(np.ones_like(correlation, dtype='bool'))
f, ax = plt.subplots(figsize=(11, 9))
#colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)
#heatmap using seaborn
sns.heatmap(correlation, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.show()
```



We are interested in correlation of features with (target variable) review_score

```
In [56]: data.corr()["review_score"]
```

```
Out[56]: payment_sequential      0.007774
payment_installments    -0.043959
payment_value           -0.083140
review_score             1.000000
zip_code_prefix_customer -0.027236
lat_customer            -0.037309
lng_customer            -0.042775
product_name_lenght     -0.013654
product_description_lenght 0.014226
product_photos_qty       0.021811
product_weight_g         -0.027676
product_length_cm        -0.020965
product_height_cm        -0.023773
product_width_cm         -0.012380
order_item_id            -0.138087
price                   0.002252
freight_value            -0.034503
zip_code_prefix_seller   0.026792
lat_seller               -0.009388
lng_seller               0.012411
Name: review_score, dtype: float64
```

It seems that none of the existing numerical features are highly correlated with the target variable.

Surprisingly price is very very less correlated with the review scoring. Similarly freight value.

2.4 Univariate analysis

```
In [57]: #Let us split the data with respect to review rating.
```

```
In [58]: rate_1 = data[data["review_score"]==1]
rate_2 = data[data["review_score"]==2]
rate_3 = data[data["review_score"]==3]
rate_4 = data[data["review_score"]==4]
rate_5 = data[data["review_score"]==5]
```

2.4.1 payment_type analysis with review_score

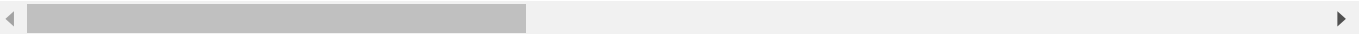
```
In [59]: #Let us check some relation in type and sequential and installments
```

```
In [60]: data[(data["payment_type"]!="voucher") & (data["payment_sequential"]>3)]
```

```
Out[60]:
```

order_id	payment_sequential	payment_type	payment_installments	payment_value	customer_id	order_status	order_purchase
----------	--------------------	--------------	----------------------	---------------	-------------	--------------	----------------

0 rows × 8 columns

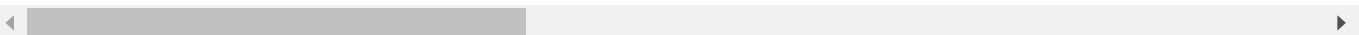


```
In [61]: data[(data["payment_sequential"]>3) & (data["payment_installments"]>1)]
```

```
Out[61]:
```

order_id	payment_sequential	payment_type	payment_installments	payment_value	customer_id	order_status	order_purchase
----------	--------------------	--------------	----------------------	---------------	-------------	--------------	----------------

0 rows × 8 columns

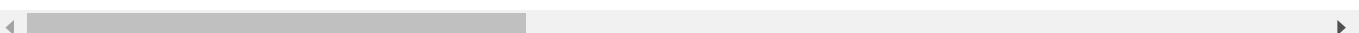


```
In [62]: data[(data["payment_type"]=="voucher") & (data["payment_installments"]>1)]
```

```
Out[62]:
```

order_id	payment_sequential	payment_type	payment_installments	payment_value	customer_id	order_status	order_purchase
----------	--------------------	--------------	----------------------	---------------	-------------	--------------	----------------

0 rows × 8 columns



- For payment sequential more than 3, all the payment methods are voucher.
- payment installments more than 1 present for payment_sequential 1, 2, 3. For payment_sequential more than 3, there is only one installment.
- For all voucher payment methods, installment is 1.

In [66]: `print(data["payment_type"].value_counts(normalize=True))`

```
credit_card    0.737854
boleto         0.194580
voucher        0.053225
debit_card     0.014341
Name: payment_type, dtype: float64
```

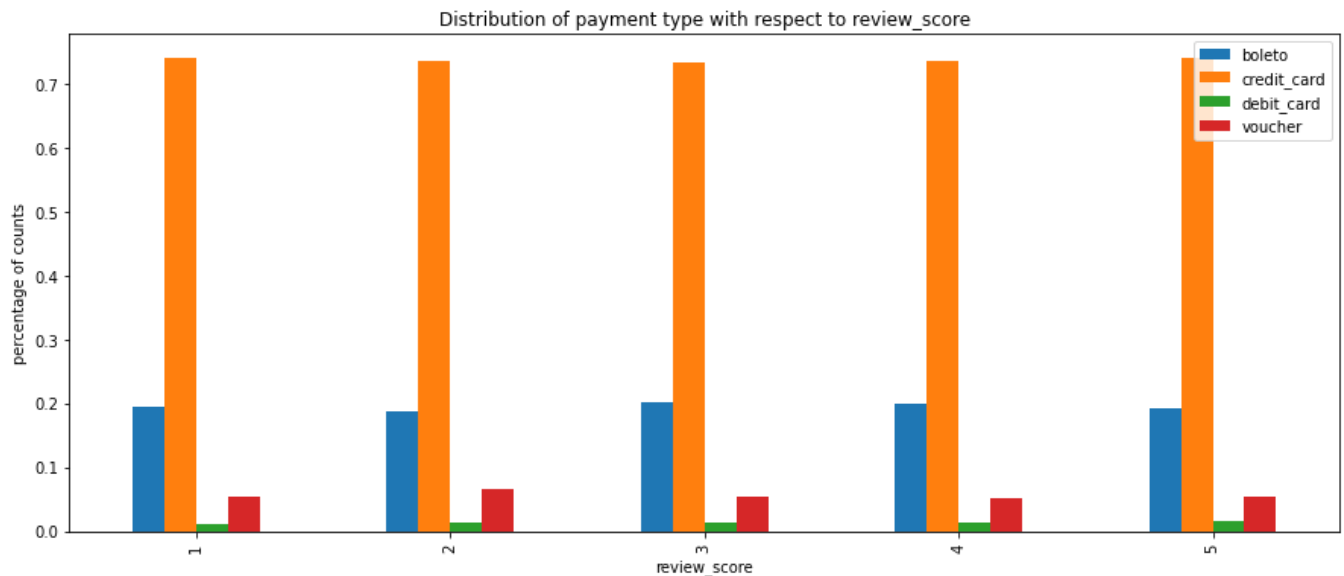
In [67]: `#payment_type grouped by review_score`
`pay_type_score = data.groupby("review_score")["payment_type"].value_counts(normalize=True)`
`pay_type_score.unstack()`

Out[67]:

	payment_type	boleto	credit_card	debit_card	voucher
	review_score				
1		0.195227	0.740427	0.011253	0.053093
2		0.186561	0.734284	0.013489	0.065666
3		0.201922	0.731954	0.012744	0.053379
4		0.198113	0.736110	0.013763	0.052014
5		0.192661	0.739004	0.015456	0.052880

Distribution of different payment types in each review rating is almost same as in the original data. Each of the ratings has nearly 74% of credit card payment. Roughly 1.2%-1.4% debit card payment.

In [68]: `pay_type_score.unstack().plot.bar(figsize=(15,6)).legend(loc="best")`
`plt.ylabel("percentage of counts")`
`plt.title("Distribution of payment type with respect to review_score")`
`plt.show()`



From the plot also we can see that, the distribution is same

Let us test statistically, whether payment_type and review_score dependent or independent.

Chi-square test with significance level alpha=0.05

Null hypothesis : payment_type and review_score are not related i.e. independent
Althernative hypothesis: payment_type and review_score are related i.e. dependent

```
In [69]: #reference: https://machinelearningmastery.com/chi-squared-test-for-machine-learning/
from scipy.stats import chi2_contingency,chi2

#creating contengency table
table = pd.crosstab(data["payment_type"],data["review_score"],margins=False)

#chi_square test using scipy.stats Library
chi_2,p_value,dof,expected = chi2_contingency(table)
```

```
In [70]: alpha = 0.05
print("Level of significance : ",alpha)
print("p-value is : ",p_value)

if p_value < alpha:
    print("Reject null hypothesis")
else:
    print("Failed to reject null hypothesis")
print("***30)
#interpreting test statistic
prob=0.95
critical = chi2.ppf(prob, dof)

print("Critical value is : ",critical)
print("chi2(test statistic) value is : ",chi_2)

if chi_2>=critical:
    print("Reject null hypothesis")
else:
    print("Failed to reject null hypothesis")
```

```
Level of significance :  0.05
p-value is :  0.00027649502286312483
Reject null hypothesis
*****
Critical value is :  21.02606981748307
chi2(test statistic) value is :  36.42753432616247
Reject null hypothesis
```


Hence we can conclude that dependence of review_score on payment_type is statistically significant

In []:

In []:

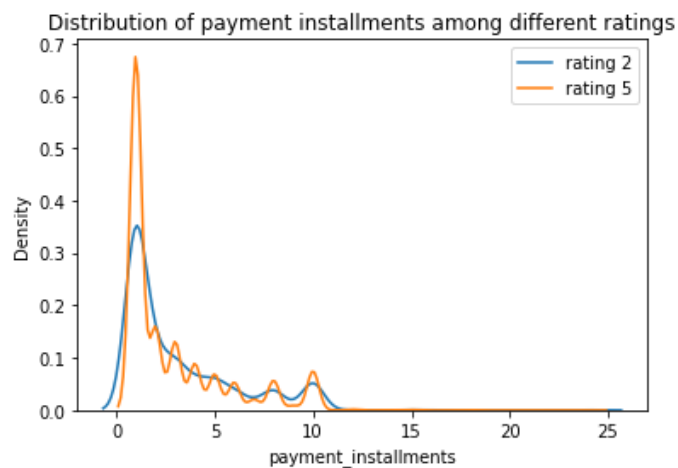
2.4.2 payment_installments analysis with review_score


```
In [71]: #plotting density plot of installments wrt each ratings
sns.FacetGrid(data,hue="review_score",height=6)\
    .map(sns.kdeplot,"payment_installments")\
    .add_legend()
plt.show()
```



```
In [72]: #distribution between rating 2 and rating 5
sns.kdeplot(rate_2["payment_installments"],label="rating 2")
sns.kdeplot(rate_5["payment_installments"],label="rating 5")


plt.title("Distribution of payment installments among different ratings")
plt.legend()
plt.show()
```



- Distributions are not separable.
- Density is peak at installments 1 for rating 5 compared to rating 2.

```
In [73]: #distribution between rating 4 and rating 5
sns.kdeplot(rate_4["payment_installments"],label="rating 4")
sns.kdeplot(rate_5["payment_installments"],label="rating 5")

plt.title("Distribution of payment installments among different ratings")
plt.legend()
plt.show()
```

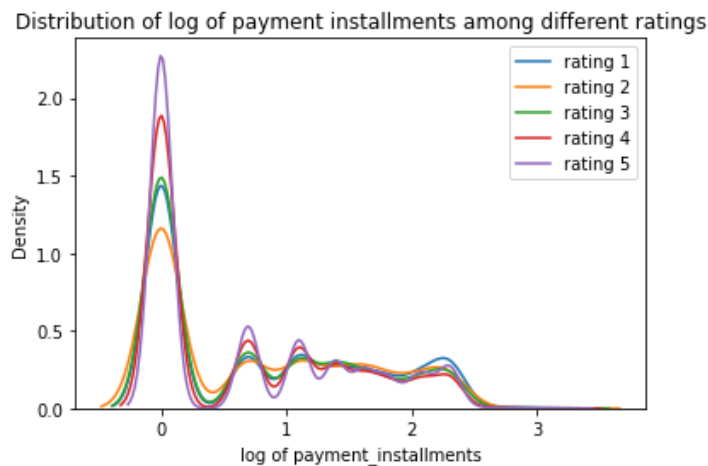


- Installment distribution of rating 4 is highly overlapping with the same of rating 5
- From the density distribution of installments we can see that, ratings are not separable using installments alone.
- Distribution plots are highly overlapped.

Since this distributon looks like follow power law, let us consider log scale.

```
In [74]: sns.kdeplot(np.log(rate_1["payment_installments"]),label="rating 1")
sns.kdeplot(np.log(rate_2["payment_installments"]),label="rating 2")
sns.kdeplot(np.log(rate_3["payment_installments"]),label="rating 3")
sns.kdeplot(np.log(rate_4["payment_installments"]),label="rating 4")
sns.kdeplot(np.log(rate_5["payment_installments"]),label="rating 5")

plt.title("Distribution of log of payment installments among different ratings")
plt.xlabel("log of payment_installments")
plt.legend()
plt.show()
```

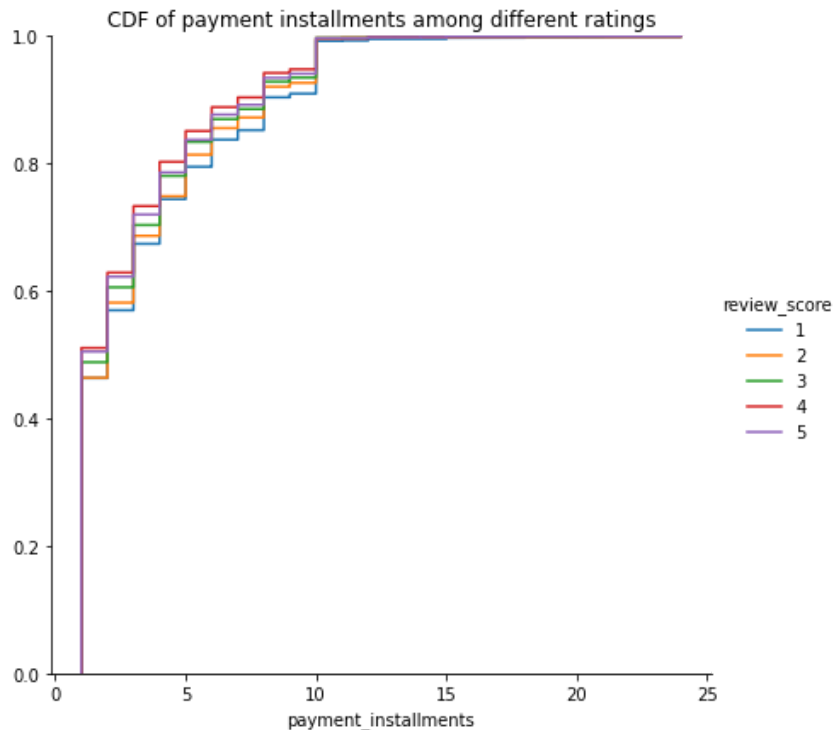


In []:

- Most of the installments = 1 in all ratings.
- Distribution of installments is almost similar among all the ratings.
- payment_installment alone cannot distinguish among the ratings.


```
In [75]: #plotting cdf
sns.FacetGrid(data,hue="review_score",height=6)\
    .map(sns.ecdfplot,"payment_installments")\
    .add_legend()

plt.title("CDF of payment installments among different ratings")
plt.xlabel("payment_installments")
plt.show()
```



- We cannot say anything more from cdf also.
- Plots are overlapped. payment_installment is discrete value. At each value the percentage of data lies below is almost same for all ratings.

```
In [76]: #boxplot
sns.boxplot(y="payment_installments",x="review_score",data=data)
plt.show()
```



```
In [77]: #mean installments
data.groupby("review_score")["payment_installments"].mean()
```

```
Out[77]: review_score
1      3.292425
2      3.149147
3      2.982451
4      2.807089
5      2.898294
Name: payment_installments, dtype: float64
```

```
In [78]: #99.9% installments
data.groupby("review_score")["payment_installments"].apply(lambda x: np.percentile(x,99.9))
```


```
Out[78]: review_score
1      24.00
2      15.36
3      15.00
4      15.00
5      15.00
Name: payment_installments, dtype: float64
```

- Obviously the first observation is box plot of all ratings overlapped.
- There are outliers in the payment_installments.
- because, from the boxplot we can see that 75% of data lies below payment_installments less than 5.
- There are very very less datapoints which have installments greater than 10 in each ratings.
- Median of rating 1,2,3 are almost same that is equal to 2. Whereas for rating 4 and 5, median is 1.
- Most of the datapoints (75%) which has rating 3, 4 or 5, have installments below 4. The same for rating 2 and 1, is 5.
- Average payment installments for rating 1 and 2 are slightly high compared to rating 3,4,5. But the difference is not significant.
- 99.9% of data points are below installments = 15 in case of ratings 2,3,4,5. Where is 99.9 percentile of rating 1 is 24.

```
In [ ]:
```

2.4.3 payment_sequential analysis with review_score

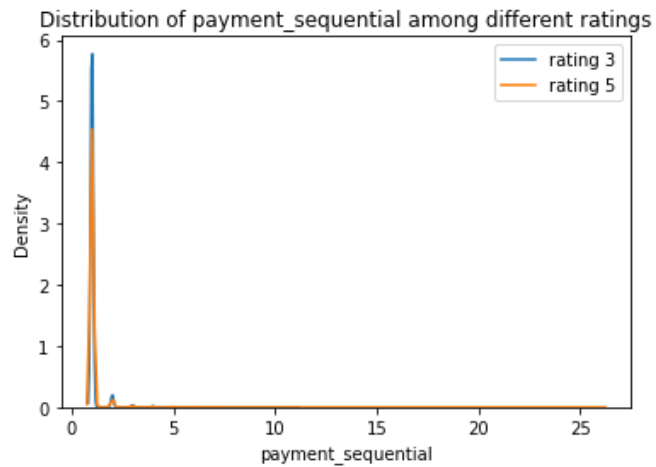
```
In [79]: #plotting density plot of installments wrt each ratings
sns.FacetGrid(data,hue="review_score",height=6)\
    .map(sns.kdeplot,"payment_sequential")\
    .add_legend()
plt.show()
```



it seems that density of payment_sequential of rating 3 is high at payment_sequential = 1 compared to other ratings. But let us check this

```
In [80]: sns.kdeplot(rate_3["payment_sequential"],label="rating 3")
sns.kdeplot(rate_5["payment_sequential"],label="rating 5")

plt.title("Distribution of payment_sequential among different ratings")
plt.xlabel("payment_sequential")
plt.legend()
plt.show()
```



Yes!! Density of rating 3 at payment_sequential=1 is peaked compared to the same of rating 5.


Distribution of payment_sequential is highly overlapped in all the ratings data. From this we can clearly see that we cannot distinguish among the ratings.

Here also we cannot get any helpful information from payment_sequential alone

```
In [81]: sns.boxplot(y="payment_sequential",x="review_score",data=data)
plt.show()

#mean installments
print("Mean \n",data.groupby("review_score")["payment_sequential"].mean())

#99.9% installments
print("99.9 percentile \n",data.groupby("review_score")["payment_sequential"].apply(lambda x: np.percentile(x,95)))
```




```
Mean
review_score
1    1.085945
2    1.095444
3    1.062572
4    1.085432
5    1.096787
Name: payment_sequential, dtype: float64
99.9 percentile
review_score
1    1.0
2    2.0
3    1.0
4    1.0
5    1.0
Name: payment_sequential, dtype: float64
```

- `payment_sequential` is also discrete feature. Its value is ranging between 1 to 26.
- From the boxplot we can see that, there are very less values of sequential greater than 1 in each rating. So these are considered as outliers in the boxplot.
- Mean value for each rating is almost same. There is no difference among the mean values.
- From 99.9 percentile value, we can clearly say that, 99.9% of data points contains sequential = 1 for all ratings except for rating 2(sequential is 2)
- It is clear that `payment_sequential` alone cannot give any information to distinguish among ratings. That means it does not have significant effect on `review_rating`.

2.4.4 payment_value analysis with review_score

```
In [82]: #plotting density plot of installments wrt each ratings
sns.FacetGrid(data,hue="review_score",height=6)\
    .map(sns.kdeplot,"payment_value")\
    .add_legend()
plt.show()
```




```
In [83]: #above density plot is overlapped. So we cannot get much insight from this. Let us have a Look at numbers.
```

```
In [ ]:
```

```
In [84]: sns.boxplot(y="payment_value",x="review_score",data=data)
plt.show()

#mean payment_value
print("Mean \n",data.groupby("review_score")["payment_value"].mean())

#90% payment_value
print("90 percentile \n",data.groupby("review_score")["payment_value"].apply(lambda x: np.percentile(x, 90)))
```



```
Mean
review_score
1    238.378781
2    185.442909
3    168.537174
4    164.140398
5    161.243232
Name: payment_value, dtype: float64
99.9 percentile
review_score
1    464.890
2    372.840
3    334.414
4    323.900
5    317.850
Name: payment_value, dtype: float64
```

- Mean values of payment_values of rating 5,4,3 are much less compared to the same of rating 1.
- But from the boxplot we can see that there are outliers with high payment values in rating 1. So this might pull the mean value of rating 1 towards higher value.
- 90th percentile value is also high for rating 1 compared to other.

In [85]:

rate_1[rate_1["payment_value"]>12000]

Out[85]:

	order_id	payment_sequential	payment_type	payment_installments	payment_value	
59174	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59175	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59176	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59177	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59178	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59179	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59180	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357
59181	03caa2c082116e1d31e67e9ae3700499	1	credit_card	1	13664.08	1617b1357

8 rows × 38 columns

- Now we can see that one customer purchased same item 8 units in the sam time. The payment value is 13664.08.
- He gave 1 rating for all this product. This high value affects the mean value of payment value.

We need to check more percentile values.

```
In [86]: for i in range(95,100):
        print("{}th percentile \n {}".format(i,data.groupby("review_score")["payment_value"].apply(lambda x
        : np.percentile(x,i))))
        print(""*40)
```

```
95th percentile
review_score
1    738.570
2    557.068
3    508.304
4    483.392
5    468.000
```

```
Name: payment_value, dtype: float64
*****
```

```
96th percentile
review_score
1    854.1020
2    622.2800
3    581.2500
4    567.2632
5    545.1400
```

```
Name: payment_value, dtype: float64
*****
```

```
97th percentile
review_score
1    1014.0200
2    696.9852
3    682.8300
4    660.4096
5    637.4604
```

```
Name: payment_value, dtype: float64
*****
```

```
98th percentile
review_score
1    1342.98
2    939.72
3    835.55
4    838.97
5    783.00
```

```
Name: payment_value, dtype: float64
*****
```

```
99th percentile
review_score
1    1853.0000
2    1308.1800
3    1076.7600
4    1150.0428
5    1071.1200
```

```
Name: payment_value, dtype: float64
*****
```

- From 95 to 99 percentile, payment value of rating 5,4,3 is significantly less rating 1.
- For 95th percentile, rating 5 has less payment value compared to rating 1.

Since there are some very larger values in each rating, lets check other percentiles.

```
In [87]: lst=[50,75,80,85]      #Let us check these percentile values
for i in lst:
    print("{}th percentile \n {}".format(i,data.groupby("review_score")["payment_value"].apply(lambda x
: np.percentile(x,i))))
    print("*****40)
```

50th percentile

review_score

1 137.57

2 119.29

3 107.78

4 104.61

5 104.43

Name: payment_value, dtype: float64

75th percentile

review_score

1 241.22

2 212.32

3 188.34

4 182.29

5 180.16

Name: payment_value, dtype: float64

80th percentile

review_score

1 284.550

2 247.640

3 218.270

4 208.632

5 207.278

Name: payment_value, dtype: float64

85th percentile

review_score

1 350.220

2 297.950

3 260.140

4 249.670

5 244.806

Name: payment_value, dtype: float64

- Median payment value is also high incase of rating 1 compared to other ratings.

From these observations, We can say, payment value has effect on review_score.

Customers who bought products of very high values most likely to give low ratings.

But this does not prove causation.

This is the insights we got from the above numbers.